

✓ Transfer Learning using ResNet101V2

Gayanthika Shankar

School of Computing and Data Science

gayanthika.s-26@scds.saiuniversity.edu.in

```
#import file
```

```
import zipfile
from google.colab import drive
drive.mount('/content/gdrive')
```

```
↗ Mounted at /content/gdrive
```

```
#Importing necessary libraries and setting random seed
import tensorflow as tf
print(tf.__version__)
```

```
from tensorflow import keras
tf.random.set_seed(42)
```

```
import numpy as np
np.random.seed(42)
```

```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
import glob
import PIL
from PIL import Image
```

```
↗ 2.17.1
```

```
from tqdm import tqdm
zip_ref = zipfile.ZipFile("/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_S
for file in tqdm(zip_ref.namelist()):
    zip_ref.extract(file, "/content/gdrive/MyDrive/DL_Data/")
zip_ref.close()
```

➞ 100%|██████████| 3332/3332 [32:18<00:00, 1.72it/s]

#Importing the images

```
imgFiles = glob.glob("/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schoon
for items in imgFiles[:8]:
    print(items)
```

➞ /content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image
/content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/image

✓ Data Preprocessing and Labelling

✓ Requirements

- Input shape: (224, 224, 3)
- 3 channels
- Width and height should be no smaller than 32
- Scale input pixels between -1 and 1

```

from tensorflow.keras.applications.resnet_v2 import preprocess_input
from tensorflow.keras.preprocessing.image import load_img, img_to_array
X = []
y = []
for fName in imgFiles:
    img = load_img(fName, target_size=(224, 224))
    img_array = img_to_array(img)
    img_preprocessed = preprocess_input(img_array)
    X.append(img_preprocessed)
    label = fName.split("/")[-2]
    y.append(label)
#Convert lists to numpy arrays
X = np.array(X)
y = np.array(y)

```

```

#Check the first few entries and their type
print(f"Type of imgFiles: {type(imgFiles)}")
print(f"Length of imgFiles: {len(imgFiles)}")
print("First few entries:")
for f in list(imgFiles)[:3]:
    print(f"- {f}, type: {type(f)}")

```

```

↔ Type of imgFiles: <class 'list'>
Length of imgFiles: 1661
First few entries:
- /content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/im
- /content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/im
- /content/gdrive/MyDrive/DL_Data/Airplanes_Motorbikes_Schooners/airplanes/im

```

```
print(y)
```

```

↔ ['airplanes' 'airplanes' 'airplanes' ... 'Motorbikes' 'Motorbikes'
'Motorbikes']

```

```
class_counts = dict()

#Count images for each class
for file_path in imgFiles:
    class_name = file_path.split("/")[-2]
    if class_name not in class_counts:
        class_counts[class_name] = 1
    else:
        class_counts[class_name] += 1

for class_name, count in class_counts.items():
    print(f"Class: {class_name}, Count: {count}")
```

```
↵ Class: airplanes, Count: 800
   Class: schooner, Count: 63
   Class: Motorbikes, Count: 798
```

```
from sklearn.preprocessing import LabelEncoder
lEncoder = LabelEncoder()
y = lEncoder.fit_transform(y)

print(set(y))
print(lEncoder.classes_)
```

```
↵ {0, 1, 2}
   ['Motorbikes' 'airplanes' 'schooner']
```

```
X = np.array(X)
y = np.array(y)

print(X.shape)
print(y.shape)
```

```
↵ (1661, 224, 224, 3)
   (1661,)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
                                                    stratify=y, random_state=42)
```

```
print("X_train_shape: {}".format(X_train.shape))
```

```
print("X_test_shape: {}".format(X_test.shape))
```

```
⇒ X_train_shape: (1245, 224, 224, 3)
   X_test_shape: (416, 224, 224, 3)
```

```
mu = X_train.mean()
```

```
std = X_train.std()
```

```
X_train_std = (X_train-mu)/std
```

```
X_test_std = (X_test-mu)/std
```

```
X_train_std.shape
```

```
⇒ (1245, 224, 224, 3)
```

```
y_train.shape
```

```
⇒ (1245,)
```

```
X_train
```

```
⇒ array([[[[ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            ...,
            [-0.36470586, -0.30196077, -0.21568626],
            [-0.3960784 , -0.34117645, -0.29411763],
            [ 0.9607843 ,  1.          ,  1.          ]]],
          [[ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            ...,
            [-0.36470586, -0.30196077, -0.21568626],
            [-0.3960784 , -0.34117645, -0.29411763],
            [ 0.9607843 ,  1.          ,  1.          ]]],
          ...,
          [[ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            [ 1.          ,  1.          ,  1.          ],
            ...,
            [-0.36470586, -0.30196077, -0.21568626],
            [-0.3960784 , -0.34117645, -0.29411763],
            [ 0.9607843 ,  1.          ,  1.          ]]]])
```

```
--
```

```
--
```

```

[[ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 ...,
 [-0.36470586, -0.30196077, -0.21568626],
 [-0.3960784 , -0.34117645, -0.29411763],
 [ 0.9607843 ,  1.          ,  1.          ]],

...,

[[ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 ...,
 [-0.3960784 , -0.42745095, -0.45098037],
 [-0.41176468, -0.44313723, -0.46666664],
 [ 1.          ,  0.99215686,  0.9764706 ]],

[[ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 ...,
 [-0.3960784 , -0.42745095, -0.45098037],
 [-0.41176468, -0.44313723, -0.46666664],
 [ 1.          ,  0.99215686,  0.9764706 ]],

[[ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 ...,
 [-0.38823527, -0.41960782, -0.44313723],
 [-0.40392154, -0.4352941 , -0.4588235 ],
 [ 1.          ,  1.          ,  0.9843137 ]]],

[[[ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 ...,
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ],
 [ 1.          ,  1.          ,  1.          ]]],

[[ 1.          ,  1.          ,  1.          ,  1.          ]

```

```
#Load a single image and check its values before preprocessing
img = load_img(imgFiles[0], target_size=(224, 224))
img_array = img_to_array(img)
print("Original image values range:", img_array.min(), "to", img_array.max())
print("Sample of original image values:\n", img_array[0:2, 0:2])

#Check after preprocessing
img_preprocessed = preprocess_input(img_array)
print("\nPreprocessed image values range:", img_preprocessed.min(), "to", img_pre|
print("Sample of preprocessed values:\n", img_preprocessed[0:2, 0:2])
```

```
⇒ Original image values range: 0.0 to 255.0
Sample of original image values:
[[[255. 255. 255.]
  [255. 255. 255.]
```

```
  [[255. 255. 255.]
   [255. 255. 255.]]]
```

```
Preprocessed image values range: -1.0 to 1.0
Sample of preprocessed values:
[[[1. 1. 1.]
  [1. 1. 1.]
```

```
  [[1. 1. 1.]
   [1. 1. 1.]]]
```


```
print("X_train shape:", X_train.shape)
print("Value range:", X_train.min(), "to", X_train.max())
```

```
⇒ X_train shape: (1245, 224, 224, 3)
Value range: -1.0 to 1.0
```

```
import matplotlib.pyplot as plt

def show_images(X, n=5):
    plt.figure(figsize=(15, 3))
    for i in range(n):
        plt.subplot(1, n, i+1)
        #Convert back from preprocessed form for visualization
        img = (X[i] + 1) / 2 #Scale range
        plt.imshow(img)
        plt.axis('off')
    plt.show()
```

```
show_images(X_train)
```



✓ Building the model


```
import tensorflow as tf
print(tf.__version__)
```

```
from tensorflow import keras
tf.random.set_seed(42)
```

```
base_model2 = keras.applications.ResNet101V2(
    include_top=True,
    weights="imagenet",
    input_tensor=None,
    input_shape=None,
    pooling=None,
    classes=1000,
    classifier_activation="softmax",
    name="resnet101v2",
)
```

2.17.1
 Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/179518384/179518384> **8s** 0us/step

```
base_model2 = keras.applications.ResNet101V2(include_top=False, input_shape = (224
```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/171317808/171317808> **8s** 0us/step

```
base_model2.trainable = False
```

```
for layer in base_model2.layers:
    layer.trainable = False
```

```
from tensorflow.keras import layers
```

```
#classifier
```

```
x = keras.layers.GlobalAveragePooling2D()(base_model2.output)
x = layers.BatchNormalization()(x)
x = keras.layers.Dropout(0.35)(x)
output_ = layers.Dense(3, activation='softmax')(x)
```

```
model2 = keras.models.Model(inputs=[base_model2.input], outputs=[output_])
```

✓ Training the model

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
import numpy as np

model2.compile(loss='sparse_categorical_crossentropy',
               optimizer='adam',
               metrics=['accuracy'])

callbacks2 = [keras.callbacks.ModelCheckpoint("best_ResNet101V2_TL.weights.h5",
                                             monitor='val_accuracy',
                                             save_weights_only=True,
                                             save_best_only=True)]

datagen = ImageDataGenerator(
    rotation_range=20, #randomly rotate images by up to 20 degrees
    width_shift_range=0.2, #randomly shift images horizontally by up to 20%
    height_shift_range=0.2, #randomly shift images vertically by up to 20%
    horizontal_flip=True, #randomly flip images horizontally
    fill_mode='nearest', #strategy for filling in newly created pixels
    validation_split=0.1 #10% validation split defined here
)

#Create train generator
train_generator = datagen.flow(
    X_train_std,
    y_train,
    batch_size=32,
    subset='training' #Specify this is for training
)

#Create validation generator
validation_generator = datagen.flow(
    X_train_std,
    y_train,
    batch_size=32,
    subset='validation' #Specify this is for validation
)











#Train the model
history = model2.fit(
    train_generator,
```

```

validation_data=validation_generator,
epochs=10,
callbacks=callbacks2
)

```

```

⇒ Epoch 1/10
  /usr/local/lib/python3.10/dist-packages/keras/src/trainers/data_adapters/py_d
    self._warn_if_super_not_called()
36/36  44s 777ms/step - accuracy: 0.9659 - loss: 0.0872 -
Epoch 2/10
36/36  13s 304ms/step - accuracy: 0.9803 - loss: 0.0551 -
Epoch 3/10
36/36  12s 283ms/step - accuracy: 0.9748 - loss: 0.0575 -
Epoch 4/10
36/36  13s 312ms/step - accuracy: 0.9799 - loss: 0.0522 -
Epoch 5/10
36/36  12s 279ms/step - accuracy: 0.9718 - loss: 0.0731 -
Epoch 6/10
36/36  12s 279ms/step - accuracy: 0.9751 - loss: 0.0577 -
Epoch 7/10
36/36  12s 282ms/step - accuracy: 0.9640 - loss: 0.0881 -
Epoch 8/10
36/36  12s 290ms/step - accuracy: 0.9765 - loss: 0.0649 -
Epoch 9/10
36/36  12s 288ms/step - accuracy: 0.9865 - loss: 0.0381 -
Epoch 10/10
36/36  12s 282ms/step - accuracy: 0.9738 - loss: 0.0727 -

```


```
#visualize accuracy
```

```
keys = ['accuracy', 'val_accuracy']  
progress = {k:v for k,v in history.history.items() if k in keys}
```

```
import pandas as pd  
pd.DataFrame(progress).plot()
```

```
plt.xlabel("epochs")  
plt.ylabel("accuracy")
```

```
plt.grid(True)  
plt.show()
```



✓ Update model with best weight



```

model2.load_weights("best_ResNet101V2_TL.weights.h5")

testLoss2, testAccuracy2 = model2.evaluate(x = X_test_std, y = y_test)

print("Test-loss: %f, Test-accuracy: %f" % (testLoss2, testAccuracy2))

```



 **13/13**  **2s** 125ms/step – accuracy: 0.9782 – loss: 0.0757
 Test-loss: 0.049212, Test-accuracy: 0.983173

✓ Performance

```

y_prob = model2.predict(X_test_std)
y_predict = np.argmax(y_prob, axis=-1)
print(y_predict)

```

 **13/13**  **6s** 125ms/step

```

[1 1 0 0 0 0 0 1 0 0 0 0 0 0 1 1 1 1 0 1 1 1 1 1 0 0 0 1 1 0 1 0 0 0 1
 0 0 0 1 0 1 0 0 1 1 0 0 1 1 1 1 0 0 1 1 1 1 0 0 0 1 0 0 0 1 1 0 0 0 1 0
 0 1 1 0 1 0 1 0 2 1 0 1 0 0 0 0 1 1 1 0 0 2 0 0 1 0 0 0 0 0 1 1 1 1 1 0 0
 0 0 0 1 0 1 1 1 1 1 0 0 0 1 1 1 1 0 2 0 1 1 1 0 1 0 0 1 1 0 1 1 1 0 1 1 1
 1 0 1 0 1 1 1 1 0 2 1 0 0 1 0 1 1 0 0 0 1 0 0 1 1 1 0 1 0 0 0 0 0 0 0 1 1
 1 0 1 0 0 1 1 0 1 1 0 1 0 2 0 0 0 1 1 1 1 0 0 0 1 0 0 0 1 2 0 0 1 0 0 0 1
 0 1 0 0 0 0 1 0 1 1 0 1 0 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 1 0 1 1 1 0 0 1
 0 1 0 1 0 0 1 1 1 0 0 1 0 0 1 1 1 0 0 0 0 1 0 1 0 1 0 0 1 0 2 2 1 1 1 1 0
 0 2 1 0 0 1 0 1 1 0 0 1 0 1 1 1 0 0 1 1 0 2 1 2 0 0 0 0 1 0 1 0 0 0 1 1 2
 1 0 1 1 1 1 0 0 1 0 1 2 0 2 0 1 1 0 0 0 0 1 1 0 0 0 1 0 0 1 0 1 1 1 1 1 1
 0 0 1 1 0 0 1 0 1 1 1 0 0 0 0 0 1 1 0 1 1 0 2 1 1 1 1 1 1 1 0 1 0 1 1 0 1
 0 1 0 2 0 0 1 0 0]

```

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

cm = confusion_matrix(y_true = y_test, y_pred = y_predict)

fig, ax = plt.subplots(figsize=(6, 6))
ax.matshow(cm, cmap=plt.cm.Purples, alpha=0.3)

for i in range(cm.shape[0]):
    for j in range(cm.shape[1]):
        ax.text(x=j, y=i, s=cm[i, j], va='center', ha='center')

```


```
ax.title.set_text('Model2 TL CF\n')
plt.xlabel('Predicted label')
plt.ylabel('True label')

plt.tight_layout()
plt.savefig("ConfusionMatrix.png", dpi=300, format='png', pad_inches=0.3)
plt.show()

print(set(y))
print(lEncoder.classes_)
```



Model2 TL CF



```
{0, 1, 2}  
['Motorbikes' 'airplanes' 'schooner']
```

```

from sklearn.metrics import precision_score, recall_score, f1_score

pScore = precision_score(y_true= y_test, y_pred = y_predict, average = 'weighted')
print("Precision: ", pScore)

rScore = recall_score(y_true= y_test, y_pred = y_predict, average = 'weighted')
print("Recall: ", rScore)

fScore = f1_score(y_true= y_test, y_pred = y_predict, average = 'weighted')
print("F1-score: ", fScore)

print("\n\n\n")

```

```

⇒ Precision: 0.9837421033073207
Recall: 0.9831730769230769
F1-score: 0.9831679220992199

```

✓ Save the model and the dataset

```
model2.save('/content/gdrive/MyDrive/DL/ResNet101V2_Best_Model_TL.h5')
```

```
⇒ WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `
```

```
from numpy import save
```

```

save('/content/gdrive/MyDrive/DL/X_train_std_model2.npy', X_train_std)
save('/content/gdrive/MyDrive/DL/X_test_std_model2.npy', X_test_std)

```

```

save('/content/gdrive/MyDrive/DL/y_train_model2.npy', y_train)
save('/content/gdrive/MyDrive/DL/y_test_model2.npy', y_test)

```